

# Introduction to Data Science

## CS61

June 12 - July 12, 2018



Dr. Ash Pahwa

---

Lesson 7: Classification - kNN

Lesson 7.1: kNN Modeling Method

k Nearest Neighbor



# Outline

---


- Similarity Based Learning
- The Bayes Classifier
- 'k' Nearest Neighbor (kNN) Model
- kNN Model Assessment
- Data Normalization:  
Standardization & Scaling
- kNN in R
  - Example 1: Dataset = Food, No Package
  - Example 2: Dataset = Food, Package = Class
  - Example 3: Dataset = Iris, Package = Class
- kNN in Python
  - Example 4: Dataset = Food, No Package
  - Example 5: Dataset = Food, Package = Scikit-Learn



# Similarity Based Learning

---

# Classification




## Characteristics:

- Color: Red, Green, Yellow
- Shape: Circle, Pentagon, Star

A new object is given to us:

- **Determine which class does it belong to?**
- **Compute the boundaries between classes**

# Similarity Based Learning




# Similarity Based Learning


- Compute the distance matrices between objects

$$\text{Euclidean distance} = \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$$


$$\text{Manhattan distance} = |x_2 - x_1| + |y_2 - y_1|$$



Manhattan



Euclidean






# The Bayes Classifier

---

Based on “Thomas Bayes” Theorem

# Classification




- Assumption#1: 2 classes  $y = \{0,1\}$
- Assumption#2: Data distribution is Gaussian (Normal Distribution)
- Compute the following
  - Given the point  $x_0$  what is the probability that belong to class 0
    - $P(y = 0|x = x_0)$
  - Given the point  $x_0$  what is the probability that belong to class 1
    - $P(y = 1|x = x_0)$
  - $P(y = 0|x = x_0) + P(y = 1|x = x_0) = 1$
- If  $P(y = 0|x = x_0) > P(y = 1|x = x_0)$ 
  - The new object belongs to class 0
- If  $P(y = 0|x = x_0) < P(y = 1|x = x_0)$ 
  - The new object belongs to class 1



# The Bayes Classifier Decision Boundary

- At decision boundary both the probabilities would be the same
- $P(y = 0|x = x_0) = P(y = 1|x = x_0)$





# k Nearest Neighbor (kNN)

---



# kNN Model

---

- Parametric or Non-Parametric
  - Non-Parametric Model
    - No assumption is made about the model parameters when we start
- Flexibility vs Ability of Interpret Results
  - Flexible = Function of 'k'
    - When the value of 'k' is low = more flexible
    - When the value of 'k' is high = less flexible
  - Ability to interpret results = low
- Supervised or Unsupervised
  - Supervised Model
    - Response variable is needed
- Regression vs Classification
  - Classification




# Pros and Cons of kNN

---

Pros	Cons
Simple and Effective	Does not produce a model, limiting the ability to understand how the features are related to the class
Makes no assumption about the underlying data distribution	Requires selection of an appropriate value of 'k'
Non-parametric	

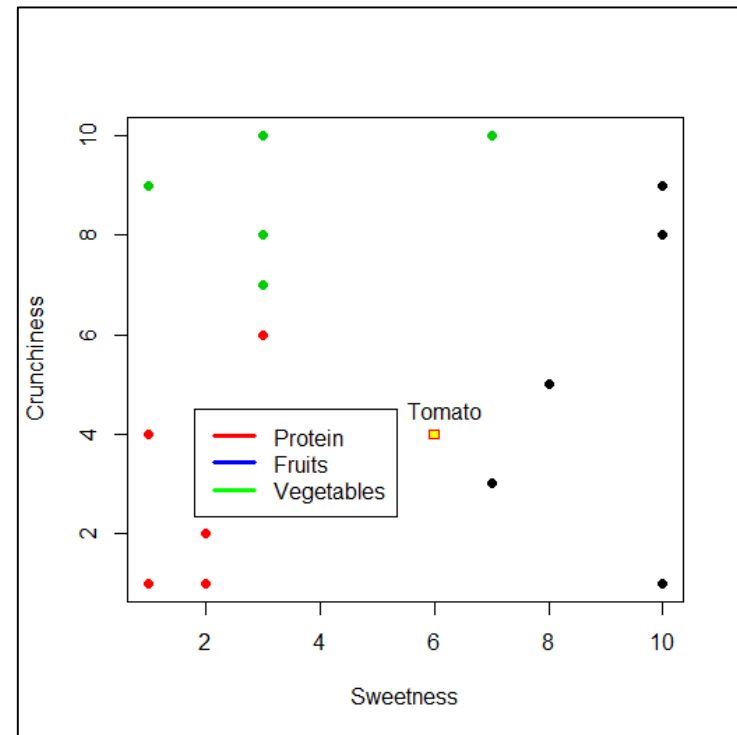

# Example

#	Item	Sweetness	Crunchiness	Food Type
1	Apple	10	9	Fruit
2	Bacon	1	4	Protein
3	Banana	10	1	Fruit
4	Carrot	7	10	Vegetable
5	Celery	3	10	Vegetable
6	Cheese	1	1	Protein
7	Grape	8	5	Fruit
8	Green bean	3	7	Vegetable
9	Nuts	3	6	Protein
10	Orange	7	3	Fruit
11	Lettuce	1	9	Vegetable
12	Cucumber	3	8	Vegetable
13	Shrimp	2	2	Protein
14	Fish	2	1	Protein
15	Pear	10	8	Fruit



# Predict Tomato: Fruit or Vegetable?

- Tomato
  - Sweetness = 6
  - Crunchiness = 4



# Distance Function

\* Euclidean

\* Manhattan

- Test data and Training Dataset with 2 dimensions


- Test Case: Point  $p = (p_1, p_2)$
- Training Data: Point  $q_1 = (q_{11}, q_{12})$
- Training Data: Point  $q_2 = (q_{21}, q_{22})$
- ...
- Training Data: Point  $q_n = (q_{n1}, q_{n2})$

- **Euclidean Distance** between  $p$  and  $q_1 = \sqrt{(p_1 - q_{11})^2 + (p_2 - q_{12})^2}$

-----

- Distance formula with k dimensions

- **Euclidean Distance** between  $p$  and  $q_1 = \sqrt{\sum_1^k (p_i - q_{1i})^2}$
- **Manhattan Distance** between  $p$  and  $q_1 = \sum_1^k |p_i - q_{1i}|$



# Euclidean Distance from Tomato

Distance between Apple and Tomato

$$\begin{aligned} \text{Euclidean distance} &= \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2} \\ &= \sqrt{(10 - 6)^2 + (9 - 4)^2} = \sqrt{4^2 + 5^2} = \sqrt{41} = 6.4 \end{aligned}$$

	A	B	C	D	E	F	G
1							
2		<b>Training Data</b>					
3		<b>#</b>	<b>Item</b>	<b>Sweetness</b>	<b>Crunchiness</b>	<b>Food Type</b>	<b>Euclidean Distance From Tomato</b>
4		1	Apple	10	9	Fruit	6.4
5		2	Bacon	1	4	Protein	5.0
6		3	Banana	10	1	Fruit	5.0
7		4	Carrot	7	10	Vegetable	6.1
8		5	Celery	3	10	Vegetable	6.7
9		6	Cheese	1	1	Protein	5.8
10		7	Grape	8	5	Fruit	2.2
11		8	Green bean	3	7	Vegetable	4.2
12		9	Nuts	3	6	Protein	3.6
13		10	Orange	7	3	Fruit	1.4
14		11	Lettuce	1	9	Vegetable	7.1
15		12	Cucumber	3	8	Vegetable	5.0
16		13	Shrimp	2	2	Protein	4.5
17		14	Fish	2	1	Protein	5.0
18		15	Pear	10	8	Fruit	5.7
19							
20		<b>Test Data</b>					
21			Tomato	6	4		
22							



# Sorted Distance

	A	B	C	D	E	F	G
1							
2		<b>Training Data</b>					
3		<b>#</b>	<b>Item</b>	<b>Sweetness</b>	<b>Crunchiness</b>	<b>Food Type</b>	<b>Euclidean Distance From Tomato</b>
4		10	Orange	7	3	Fruit	1.4
5		7	Grape	8	5	Fruit	2.2
6		9	Nuts	3	6	Protein	3.6
7		8	Green bean	3	7	Vegetable	4.2
8		13	Shrimp	2	2	Protein	4.5
9		2	Bacon	1	4	Protein	5.0
10		3	Banana	10	1	Fruit	5.0
11		12	Cucumber	3	8	Vegetable	5.0
12		14	Fish	2	1	Protein	5.0
13		15	Pear	10	8	Fruit	5.7
14		6	Cheese	1	1	Protein	5.8
15		4	Carrot	7	10	Vegetable	6.1
16		1	Apple	10	9	Fruit	6.4
17		5	Celery	3	10	Vegetable	6.7
18		11	Lettuce	1	9	Vegetable	7.1
19							

K=1

Pick the top 1 entry

Votes: Fruit = 1

Final Result = Tomato is a Fruit

	A	B	C	D	E	F	G
1							
2	Training Data						
3		#	Item	Sweetness	Crunchiness	Food Type	Euclidean Distance From Tomato
4		10	Orange	7	3	Fruit	1.4
5		7	Grape	8	5	Fruit	2.2
6		9	Nuts	3	6	Protein	3.6
7		8	Green bean	3	7	Vegetable	4.2
8		13	Shrimp	2	2	Protein	4.5
9		2	Bacon	1	4	Protein	5.0
10		3	Banana	10	1	Fruit	5.0
11		12	Cucumber	3	8	Vegetable	5.0
12		14	Fish	2	1	Protein	5.0
13		15	Pear	10	8	Fruit	5.7
14		6	Cheese	1	1	Protein	5.8
15		4	Carrot	7	10	Vegetable	6.1
16		1	Apple	10	9	Fruit	6.4
17		5	Celery	3	10	Vegetable	6.7
18		11	Lettuce	1	9	Vegetable	7.1
19							

K=3

Pick the top 3 entries

Votes: Fruit = 2, Protein = 1

Final Result = Tomato is a Fruit

	A	B	C	D	E	F	G
1							
2		Training Data					
3		#	Item	Sweetness	Crunchiness	Food Type	Euclidean Distance From Tomato
4		10	Orange	7	3	Fruit	1.4
5		7	Grape	8	5	Fruit	2.2
6		9	Nuts	3	6	Protein	3.6
7		8	Green bean	3	7	Vegetable	4.2
8		13	Shrimp	2	2	Protein	4.5
9		2	Bacon	1	4	Protein	5.0
10		3	Banana	10	1	Fruit	5.0
11		12	Cucumber	3	8	Vegetable	5.0
12		14	Fish	2	1	Protein	5.0
13		15	Pear	10	8	Fruit	5.7
14		6	Cheese	1	1	Protein	5.8
15		4	Carrot	7	10	Vegetable	6.1
16		1	Apple	10	9	Fruit	6.4
17		5	Celery	3	10	Vegetable	6.7
18		11	Lettuce	1	9	Vegetable	7.1
19							

K=9

Pick the top 9 entries

Votes: Fruit = 3, Protein = 4, Vegetable = 2

Final Result = Tomato is a Protein


	A	B	C	D	E	F	G
1							
2		Training Data					
3		#	Item	Sweetness	Crunchiness	Food Type	Euclidean Distance From Tomato
4		10	Orange	7	3	Fruit	1.4
5		7	Grape	8	5	Fruit	2.2
6		9	Nuts	3	6	Protein	3.6
7		8	Green bean	3	7	Vegetable	4.2
8		13	Shrimp	2	2	Protein	4.5
9		2	Bacon	1	4	Protein	5.0
10		3	Banana	10	1	Fruit	5.0
11		12	Cucumber	3	8	Vegetable	5.0
12		14	Fish	2	1	Protein	5.0
13		15	Pear	10	8	Fruit	5.7
14		6	Cheese	1	1	Protein	5.8
15		4	Carrot	7	10	Vegetable	6.1
16		1	Apple	10	9	Fruit	6.4
17		5	Celery	3	10	Vegetable	6.7
18		11	Lettuce	1	9	Vegetable	7.1
19							

# Results: Count the Votes

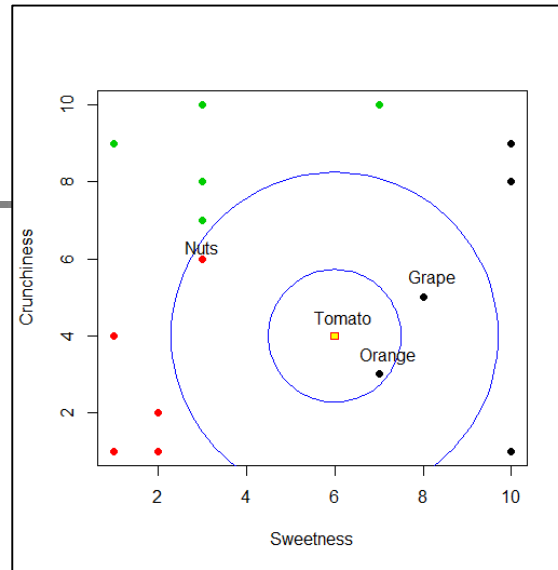
- k=1, Fruit (Fruit = 1)
- k=3, Fruit (Fruit = 2, Protein = 1)
- k=5, Fruit or Protein (Fruit = 2, Protein = 2, Vegetable = 1)
- k=9, Protein (Fruit = 3, Protein = 4, Vegetable = 2)

Training Data								
#	Item	Sweetness	Crunchiness	Food Type	Euclidean Distance From Tomato	k=1, #votes	k=3, #votes	k=5, #votes
10	Orange	7	3	Fruit	1.4	Fruit=1		
7	Grape	8	5	Fruit	2.2			
9	Nuts	3	6	Protein	3.6		Fruit=2, Protein=1	
8	Green bean	3	7	Vegetable	4.2			
13	Shrimp	2	2	Protein	4.5			Fruit =2. Protein=2, Vegetable=1
2	Bacon	1	4	Protein	5.0			
3	Banana	10	1	Fruit	5.0			
12	Cucumber	3	8	Vegetable	5.0			
14	Fish	2	1	Protein	5.0			
15	Pear	10	8	Fruit	5.7			
6	Cheese	1	1	Protein	5.8			
4	Carrot	7	10	Vegetable	6.1			
1	Apple	10	9	Fruit	6.4			
5	Celery	3	10	Vegetable	6.7			
11	Lettuce	1	9	Vegetable	7.1			

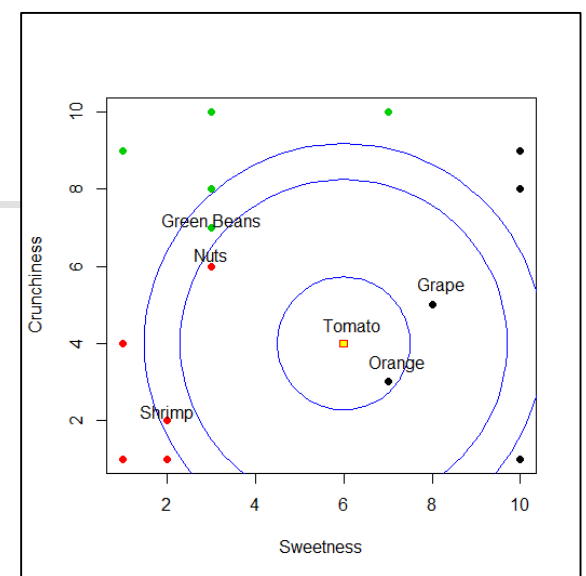
# Results: Visualization



K=1



K=3



K=5

Training Data								
#	Item	Sweetness	Crunchiness	Food Type	Euclidean Distance From Tomato	k=1, #votes	k=3, #votes	k=5, #votes
10	Orange	7	3	Fruit	1.4	Fruit=1		
7	Grape	8	5	Fruit	2.2			
9	Nuts	3	6	Protein	3.6		Fruit=2, Protein=1	
8	Green bean	3	7	Vegetable	4.2			
13	Shrimp	2	2	Protein	4.5			Fruit =2. Protein=2, Vegetable=1
2	Bacon	1	4	Protein	5.0			
3	Banana	10	1	Fruit	5.0			
12	Cucumber	3	8	Vegetable	5.0			
14	Fish	2	1	Protein	5.0			
15	Pear	10	8	Fruit	5.7			
6	Cheese	1	1	Protein	5.8			
4	Carrot	7	10	Vegetable	6.1			
1	Apple	10	9	Fruit	6.4			
5	Celery	3	10	Vegetable	6.7			
11	Lettuce	1	9	Vegetable	7.1			




# kNN Model Assessment

---

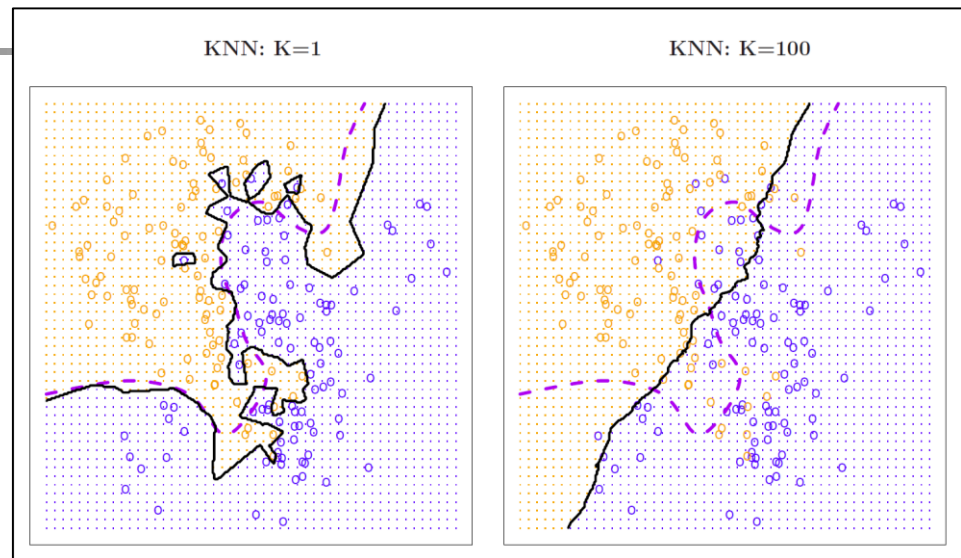
How to decide the value of 'k'?

# What Should be the Value of 'k'?


Bayes Decision Boundary: Purple dashed line



kNN: k=10: Black Line



kNN: k=1: Black Line



kNN: k=100: Black Line

- If the natural boundary is non-linear
  - Lower value of 'k' is better
- If the natural boundary is linear
  - Higher value of 'k' is better





# Bias Variance Tradeoff

---

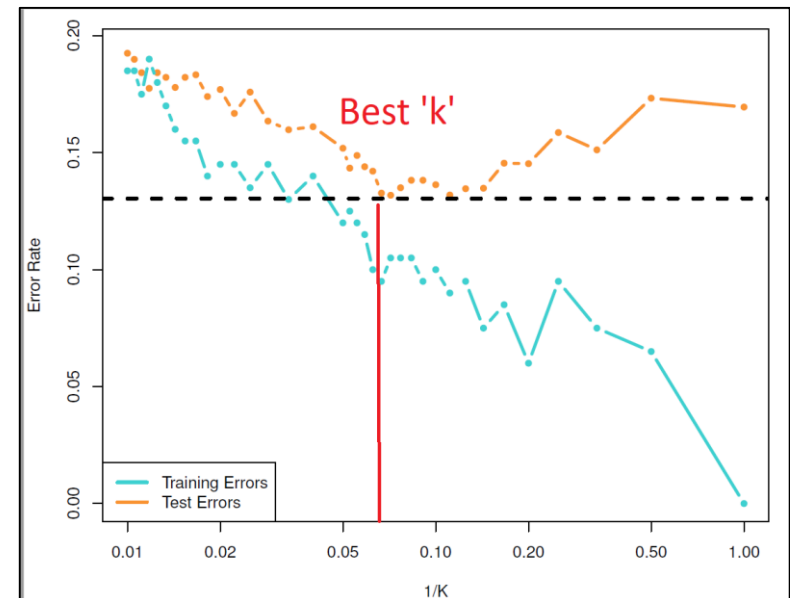

- If ' $k$ ' = 1
  - Decision boundary is overly flexible
  - Model finds patterns in the data that don't correspond to the Bayes boundary
  - Low bias but high variance
- If ' $k$ ' = 100 (large)
  - Every training instance is represented in the final vote –
  - All testing data will be classified as the class with majority votes
  - Decision boundary becomes close to linear
  - Model becomes less flexible
  - Low variance but high bias

# Testing Error as a Function of 'k'

## ■ Best 'k'

■  $1/k = 0.07$

■  $'k'=14$





# Data Normalization: Standardization & Scaling

---

In kNN Modeling we usually either  
Standardize or Scale the data



# Data Standardization & Scaling

---

- Suppose we have 2 data items
  - Height: varies from 4 – 7 feet
  - Net Worth: \$10,000 - \$100B
- If we use both the variables in a model
  - Net Worth will dominate because it contains large values
- Solution
  - Standardize
  - Scale



# Data Standardization and Scaling

---

- Standardization Data Variation
  - -3 to +3

$$z = \frac{\text{Data Value} - \text{Mean}}{\text{Standard Deviation}} = \frac{y - \mu}{\sigma}$$

- Scaling Data Variation
  - 0 to 1

$$y_i^j = \frac{x_i^j - \min_j}{\max_j - \min_j}$$

# Example

```
> normalize = function(x) {
+   return( (x-min(x))/(max(x)-min(x))) }
> data = c(124,3,311,341,298,136,23,75,5,51,822,364,663,444,999)
> (standard.data = scale(data))
      [,1]
[1,] -0.603086904
[2,] -0.994156118
[3,]  0.001292791
[4,]  0.098252100
[5,] -0.040722910
[6,] -0.564303180
[7,] -0.929516578
[8,] -0.761453776
[9,] -0.987692164
[10,] -0.839021223
[11,]  1.652833026
[12,]  0.172587571
[13,]  1.138948687
[14,]  0.431145729
[15,]  2.224892951
attr(,"scaled:center")
[1] 310.6
attr(,"scaled:scale")
[1] 309.4081
> (normalized.data = normalize(data))
[1] 0.121485944 0.000000000 0.309236948 0.339357430 0.296184739
0.133534137 0.020080321 0.072289157 0.002008032 0.048192771 0.822289157
[12] 0.362449799 0.662650602 0.442771084 1.000000000
>
```

	A	B	C	D	E	F	
1		#	Data	Standardization		Scaling	
2		1	124	-0.60		0.12	
3		2	3	-0.99		0.00	
4		3	311	0.00		0.31	
5		4	341	0.10		0.34	
6		5	298	-0.04		0.30	
7		6	136	-0.56		0.13	
8		7	23	-0.93		0.02	
9		8	75	-0.76		0.07	
10		9	5	-0.99		0.00	
11		10	51	-0.84		0.05	
12		11	822	1.65		0.82	
13		12	364	0.17		0.36	
14		13	663	1.14		0.66	
15		14	444	0.43		0.44	
16		15	999	2.22		1	
17							
18		Mean	310.60		Minimum	3	
19		StdDev	309.41		Maximum	999	
20							



# Building kNN Model in R

---



# Example 1

---

Dataset = Food

Method: Build the kNN algorithm using R



# Example 1: Food Dataset

## Read the Dataset

```
> # import the CSV file
> food <- read.csv("03 food16.csv")
> food
```

	X.	Item	Sweetness	Crunchiness	Food.Type
1	1	Apple	10	9	Fruit
2	2	Bacon	1	4	Protein
3	3	Banana	10	1	Fruit
4	4	Carrot	7	10	Vegetable
5	5	Celery	3	10	Vegetable
6	6	Cheese	1	1	Protein
7	7	Grape	8	5	Fruit
8	8	Green bean	3	7	Vegetable
9	9	Nuts	3	6	Protein
10	10	Orange	7	3	Fruit
11	11	Lettuce	1	9	Vegetable
12	12	Cucumber	3	8	Vegetable
13	13	Shrimp	2	2	Protein
14	14	Fish	2	1	Protein
15	15	Pear	10	8	Fruit
16	16	Tomato	6	4	

```
> table(food$Food.Type)
```

	Fruit	Protein	Vegetable
	1	5	5

```
>
```



# Example 1: Food Dataset

## Separate Train and Test Data

---

```
> trainData <- food[1:15,3:5]
> trainData
  Sweetness Crunchiness Food.Type
1         10           9    Fruit
2          1           4   Protein
3         10           1    Fruit
4          7          10 Vegetable
5          3          10 Vegetable
6          1           1   Protein
7          8           5    Fruit
8          3           7 Vegetable
9          3           6   Protein
10         7           3    Fruit
11         1           9 Vegetable
12         3           8 Vegetable
13         2           2   Protein
14         2           1   Protein
15        10           8    Fruit
> testData <- food[16:16,3:5]
> testData
  Sweetness Crunchiness Food.Type
16         6           4
```

# Example 1: Food Dataset

## Compute the Distance : Sort

```
> test.count = 1
> sum = rep(0,dim(trainData)[1])
> for ( i in 1:dim(trainData)[1] ) {
+   sum[i] = sum[i] + (trainData$Sweetness[i] - testData$Sweetness[test.count])^2
+   sum[i] = sum[i] + (trainData$Crunchiness[i] - testData$Crunchiness[test.count])^2
+ }
> sum
[1] 41 25 25 37 45 34 5 18 13 2 50 25 20 25 32
> (trainData$dist = sqrt(sum))
[1] 6.403124 5.000000 5.000000 6.082763 6.708204 5.830952 2.236068 4.242641 3.605551
1.414214 7.071068 5.000000 4.472136 5.000000 5.656854
> (trainData_Sorted = trainData[order(trainData$dist),])
  Sweetness Crunchiness Food.Type      dist
10         7          3      Fruit 1.414214
7          8          5      Fruit 2.236068
9          3          6    Protein 3.605551
8          3          7  Vegetable 4.242641
13         2          2    Protein 4.472136
2          1          4    Protein 5.000000
3         10          1      Fruit 5.000000
12         3          8  Vegetable 5.000000
14         2          1    Protein 5.000000
15        10          8      Fruit 5.656854
6          1          1    Protein 5.830952
4          7         10  Vegetable 6.082763
1         10          9      Fruit 6.403124
5          3         10  Vegetable 6.708204
11         1          9  Vegetable 7.071068
>
```

# Example 1: Food Dataset

## Retrieve the Nearest Neighbor

- k=1, Fruit (Fruit = 1)
- k=3, Fruit (Fruit = 2, Protein = 1)

```
> #####  
> k = 1  
> (nearestNeighbour = as.character(trainData_Sorted$Food.Type[1:k]))  
[1] "Fruit"  
> table(nearestNeighbour)  
nearestNeighbour  
Fruit  
  1  
> #####  
> k = 3  
> (nearestNeighbour = as.character(trainData_Sorted$Food.Type[1:k]))  
[1] "Fruit" "Fruit" "Protein"  
> table(nearestNeighbour)  
nearestNeighbour  
  Fruit Protein  
    2      1
```

# Example 1: Food Dataset

## Retrieve the Nearest Neighbor

- k=5, Fruit or Protein (Fruit = 2, Protein = 2, Vegetable = 1)
- k=9, Protein (Fruit = 3, Protein = 4, Vegetable = 2)

```
> #####  
> k = 5  
> (nearestNeighbour = as.character(trainData_Sorted$Food.Type[1:k]))  
[1] "Fruit"      "Fruit"      "Protein"    "Vegetable" "Protein"  
> table(nearestNeighbour)  
nearestNeighbour  
   Fruit   Protein Vegetable  
      2       2         1  
> #####  
> k = 9  
> (nearestNeighbour = as.character(trainData_Sorted$Food.Type[1:k]))  
[1] "Fruit"      "Fruit"      "Protein"    "Vegetable" "Protein"    "Protein"    "Fruit"  
"Vegetable" "Protein"  
> table(nearestNeighbour)  
nearestNeighbour  
   Fruit   Protein Vegetable  
      3       4         2
```



## Example 2

---

Dataset = Food  
Package: R/Class

# Example 1: Food Dataset

## Package: Class

```
> #####  
> # import the CSV file  
> food <- read.csv("03 food16.csv")  
> food  
  X.      Item Sweetness Crunchiness Food.Type  
1  1      Apple        10           9      Fruit  
2  2      Bacon         1           4    Protein  
3  3     Banana        10           1      Fruit  
4  4     Carrot         7          10  Vegetable  
5  5     Celery         3          10  Vegetable  
6  6     Cheese         1           1    Protein  
7  7      Grape         8           5      Fruit  
8  8 Green bean         3           7  Vegetable  
9  9        Nuts         3           6    Protein  
10 10     Orange         7           3      Fruit  
11 11     Lettuce         1           9  Vegetable  
12 12    Cucumber         3           8  Vegetable  
13 13     Shrimp         2           2    Protein  
14 14       Fish         2           1    Protein  
15 15       Pear        10           8      Fruit  
16 16     Tomato         6           4  
> table(food$Food.Type)  
  
      Fruit      Protein  Vegetable  
          1           5           5
```



# Example 1: Food Dataset

```
> train <- food[1:15,3:4]
> train
  Sweetness Crunchiness
1         10          9
2          1          4
3         10          1
4          7         10
5          3         10
6          1          1
7          8          5
8          3          7
9          3          6
10         7          3
11         1          9
12         3          8
13         2          2
14         2          1
15        10          8
> test <- food[16:16,3:4]
> test
  Sweetness Crunchiness
16         6          4
>
```



# Use 'Class' Library: kNN Function

- k=1, Fruit (Fruit = 1)
- k=3, Fruit (Fruit = 2, Protein = 1)
- k=5, Fruit or Protein (Fruit = 2, Protein = 2, Vegetable = 1)
- k=9, Protein (Fruit = 3, Protein = 4, Vegetable = 2)

```
> # load the "class" library
> library(class)
> #####
> test_pred <- knn(train = train, test = test,
+                  cl = train_labels, k = 1)
> test_pred
[1] Fruit
Levels: Fruit Protein Vegetable
> #####
> test_pred <- knn(train = train, test = test,
+                  cl = train_labels, k = 3)
> test_pred
[1] Fruit
Levels: Fruit Protein Vegetable
> #####
> test_pred <- knn(train = train, test = test,
+                  cl = train_labels, k = 5)
> test_pred
[1] Fruit
Levels: Fruit Protein Vegetable
> #####
> test_pred <- knn(train = train, test = test,
+                  cl = train_labels, k = 9)
> test_pred
[1] Protein
Levels:
```



# Example 3

---

Dataset = Iris  
Package: R/Class

# Iris Dataset

- Edgar Anderson's Iris Data
- Iris Species
  - Setosa, Versicolor, Virginica
- Sepal + Petal length and width in centimeters
- 150 records (50 flowers from each 3 species)

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa


Iris Virginica



Iris Versicolor



Iris Setosa



# Iris Dataset

```
> dim(iris)
[1] 150    5
> summary(iris)
  Sepal.Length    Sepal.Width    Petal.Length    Petal.Width      Species
Min.      :4.300   Min.      :2.000   Min.      :1.000   Min.      :0.100   setosa      :50
1st Qu.:5.100   1st Qu.:2.800   1st Qu.:1.600   1st Qu.:0.300   versicolor:50
Median :5.800   Median :3.000   Median :4.350   Median :1.300   virginica  :50
Mean    :5.843   Mean    :3.057   Mean    :3.758   Mean    :1.199
3rd Qu.:6.400   3rd Qu.:3.300   3rd Qu.:5.100   3rd Qu.:1.800
Max.    :7.900   Max.    :4.400   Max.    :6.900   Max.    :2.500
> head(iris)
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species
1          5.1         3.5         1.4         0.2   setosa
2          4.9         3.0         1.4         0.2   setosa
3          4.7         3.2         1.3         0.2   setosa
4          4.6         3.1         1.5         0.2   setosa
5          5.0         3.6         1.4         0.2   setosa
6          5.4         3.9         1.7         0.4   setosa
> names(iris)
[1] "Sepal.Length" "Sepal.Width"  "Petal.Length" "Petal.Width"  "Species"
```

```
> str(iris)
'data.frame':    150 obs. of  5 variables:
 $ Sepal.Length: num  5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
 $ Sepal.Width : num  3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
 $ Petal.Length: num  1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
 $ Petal.Width : num  0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
 $ Species      : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1 1
...
> table(iris$Species)

  setosa versicolor virginica 
      50         50         50 
>
```



# Shuffle the Dataset

```
> #####  
> # 1. Mix all the rows  
> # Like shuffle deck of cards  
> #  
> set.seed(9850)  
> gp = runif(nrow(iris)) # Generate 150 random numbers uniformly distributed  
> iris = iris[order(gp),]  
> head(iris,10)  
      Sepal.Length Sepal.Width Petal.Length Petal.Width   Species  
103           7.1         3.0         5.9         2.1  virginica  
20            5.1         3.8         1.5         0.3    setosa  
63            6.0         2.2         4.0         1.0  versicolor  
17            5.4         3.9         1.3         0.4    setosa  
83            5.8         2.7         3.9         1.2  versicolor  
53            6.9         3.1         4.9         1.5  versicolor  
118           7.7         3.8         6.7         2.2  virginica  
91            5.5         2.6         4.4         1.2  versicolor  
80            5.7         2.6         3.5         1.0  versicolor  
43            4.4         3.2         1.3         0.2    setosa  
>
```

## Normalize the dataset (from 0 – 1)

```
> #####  
> # 2. Normalize all numbers from 0 - 1  
> #  
> summary(iris[,c(1,2,3,4)])  
  Sepal.Length    Sepal.Width    Petal.Length    Petal.Width  
Min.   :4.300    Min.   :2.000    Min.   :1.000    Min.   :0.100  
1st Qu.:5.100    1st Qu.:2.800    1st Qu.:1.600    1st Qu.:0.300  
Median :5.800    Median :3.000    Median :4.350    Median :1.300  
Mean   :5.843    Mean   :3.057    Mean   :3.758    Mean   :1.199  
3rd Qu.:6.400    3rd Qu.:3.300    3rd Qu.:5.100    3rd Qu.:1.800  
Max.   :7.900    Max.   :4.400    Max.   :6.900    Max.   :2.500  
>  
> normalize = function(x) {  
+   return( (x-min(x))/(max(x)-min(x))) }  
>  
> iris_n = as.data.frame(lapply(iris[,c(1,2,3,4)],normalize))  
> summary(iris_n)  
  Sepal.Length    Sepal.Width    Petal.Length    Petal.Width  
Min.   :0.0000    Min.   :0.0000    Min.   :0.0000    Min.   :0.00000  
1st Qu.:0.2222    1st Qu.:0.3333    1st Qu.:0.1017    1st Qu.:0.08333  
Median :0.4167    Median :0.4167    Median :0.5678    Median :0.50000  
Mean   :0.4287    Mean   :0.4406    Mean   :0.4675    Mean   :0.45806  
3rd Qu.:0.5833    3rd Qu.:0.5417    3rd Qu.:0.6949    3rd Qu.:0.70833  
Max.   :1.0000    Max.   :1.0000    Max.   :1.0000    Max.   :1.00000  
>
```



# Build Training and Test Dataset

---

```
> #####  
> # 3. Build Training and Test dataset  
> #  
> iris_train = iris_n[1:129,]  
> iris_test = iris_n[130:150,]  
> iris_training_target = iris[1:129,5]  
> iris_test_target = iris[130:150,5]  
>
```

# Train the Model using Training Dataset, Predict Response Variable of the Test Dataset

## Package "class" has kNN module

```
> #####
> # call the kNN function
> #
> # The value of 'k' should be sqrt of observation
> # The value of 'k' should be an odd number
> # If a voting tie occurs, it can be resolved
> #
> # Observations = 150
> # k = 13
> #
> library(class)
>
> k = 13
> m1 <- knn(train=iris_train, test=iris_test, cl=iris_training_target, k=k)
> (t = table(iris_test_target,m1))
      m1
iris_test_target setosa versicolor virginica
      setosa      7           0           0
versicolor      0           3           2
virginica       0           0           9
>
> #####
> # Compute Accuracy of Prediction
> (accuracy = sum(diag(t))/sum(t)*100)
[1] 90.47619
>
```





## Example 4

---

Dataset = Food

Method: Build the kNN algorithm using  
Python

# Read Data Set

```
import numpy as np
import pandas as pd
from collections import Counter
#####
# Read Food Dataset
train = pd.read_csv('03 food1-15.csv')
test = pd.read_csv('05 food16.csv')
train
Out[9]:
```

	#	Item	Sweetness	Crunchiness	FoodType
0	1	Apple	10	9	Fruit
1	2	Bacon	1	4	Protein
2	3	Banana	10	1	Fruit
3	4	Carrot	7	10	Vegetable
4	5	Celery	3	10	Vegetable
5	6	Cheese	1	1	Protein
6	7	Grape	8	5	Fruit
7	8	Green bean	3	7	Vegetable
8	9	Nuts	3	6	Protein
9	10	Orange	7	3	Fruit
10	11	Lettuce	1	9	Vegetable
11	12	Cucumber	3	8	Vegetable
12	13	Shrimp	2	2	Protein
13	14	Fish	2	1	Protein
14	15	Pear	10	8	Fruit

```
test
Out[10]:
```

	#	Item	Sweetness	Crunchiness	FoodType
0	1	Tomato	6	4	NaN



# Compute the Distance

```
#####  
# Compute the distance  
# from Test object to all the Train's objects  
#  
trainC = train.shape[0]  
print(trainC)  
15  
sum = np.zeros(trainC)  
  
for i in range (0, trainC):  
    sum[i] = sum[i] + (train.Sweetness[i] - test.Sweetness[0])**2  
    sum[i] = sum[i] + (train.Crunchiness[i] - test.Crunchiness[0])**2  
  
distance = np.sqrt(sum)  
  
print(sum)  
[ 41.  25.  25.  37.  45.  34.   5.  18.  13.   2.  50.  25.  20.  25.  32.]  
  
print(distance)  
[ 6.40312424  5.          5.          6.08276253  6.70820393  5.83095189  
 2.23606798  4.24264069  3.60555128  1.41421356  7.07106781  5.  
 4.47213595  5.          5.65685425]
```



# Compute the Distance

---

```
train['dist'] = distance
```

```
print(train)
```

	#	Item	Sweetness	Crunchiness	FoodType	dist
0	1	Apple	10	9	Fruit	6.403124
1	2	Bacon	1	4	Protein	5.000000
2	3	Banana	10	1	Fruit	5.000000
3	4	Carrot	7	10	Vegetable	6.082763
4	5	Celery	3	10	Vegetable	6.708204
5	6	Cheese	1	1	Protein	5.830952
6	7	Grape	8	5	Fruit	2.236068
7	8	Green bean	3	7	Vegetable	4.242641
8	9	Nuts	3	6	Protein	3.605551
9	10	Orange	7	3	Fruit	1.414214
10	11	Lettuce	1	9	Vegetable	7.071068
11	12	Cucumber	3	8	Vegetable	5.000000
12	13	Shrimp	2	2	Protein	4.472136
13	14	Fish	2	1	Protein	5.000000
14	15	Pear	10	8	Fruit	5.656854



# Sort the Distances

```
#####  
# Sort the dataset by distance  
#  
trainSorted = train.sort_values(['dist'])  
  
print(trainSorted)
```

	#	Item	Sweetness	Crunchiness	FoodType	dist
9	10	Orange	7	3	Fruit	1.414214
6	7	Grape	8	5	Fruit	2.236068
8	9	Nuts	3	6	Protein	3.605551
7	8	Green bean	3	7	Vegetable	4.242641
12	13	Shrimp	2	2	Protein	4.472136
1	2	Bacon	1	4	Protein	5.000000
2	3	Banana	10	1	Fruit	5.000000
11	12	Cucumber	3	8	Vegetable	5.000000
13	14	Fish	2	1	Protein	5.000000
14	15	Pear	10	8	Fruit	5.656854
5	6	Cheese	1	1	Protein	5.830952
3	4	Carrot	7	10	Vegetable	6.082763
0	1	Apple	10	9	Fruit	6.403124
4	5	Celery	3	10	Vegetable	6.708204
10	11	Lettuce	1	9	Vegetable	7.071068

# Find the Nearest Neighbors

- k=1, Fruit (Fruit = 1)
- k=3, Fruit (Fruit = 2, Protein = 1)

```
#####  
# Find the nearest neighbor  
#  
k = 1  
nearestNeighbor = trainSorted.FoodType[0:k]  
print(nearestNeighbor)  
9      Fruit  
Name: FoodType, dtype: object  
  
Counter(nearestNeighbor)  
Out[36]: Counter({'Fruit': 1})  
  
k = 3  
nearestNeighbor = trainSorted.FoodType[0:k]  
print(nearestNeighbor)  
9      Fruit  
6      Fruit  
8      Protein  
Name: FoodType, dtype: object  
  
Counter(nearestNeighbor)  
Out[40]: Counter({'Fruit': 2, 'Protein': 1})
```



# Example 5

---

Dataset = Food

Package: Python/Scikit-Learn

# Example 4

## Read Food Dataset

```
import numpy as np
import pandas as pd
from sklearn import neighbors
```

```
df = pd.read_csv('04 food16.csv')
```

```
df
```

```
Out[5]:
```

	#	Item	Sweetness	Crunchiness	Food Type
0	1	Apple	10	9	Fruit
1	2	Bacon	1	4	Protein
2	3	Banana	10	1	Fruit
3	4	Carrot	7	10	Vegetable
4	5	Celery	3	10	Vegetable
5	6	Cheese	1	1	Protein
6	7	Grape	8	5	Fruit
7	8	Green bean	3	7	Vegetable
8	9	Nuts	3	6	Protein
9	10	Orange	7	3	Fruit
10	11	Lettuce	1	9	Vegetable
11	12	Cucumber	3	8	Vegetable
12	13	Shrimp	2	2	Protein
13	14	Fish	2	1	Protein
14	15	Pear	10	8	Fruit
15	16	Tomato	6	4	NaN



# Example 4

## Split data into Train + Test

```
#####  
# Split data into train + test  
#  
X = np.array(df[['Sweetness', 'Crunchiness']])  
y = df['Food Type']  
  
X_train = X[0:15,]  
X_train  
Out[15]:  
array([[10,  9],  
       [ 1,  4],  
       [10,  1],  
       [ 7, 10],  
       [ 3, 10],  
       [ 1,  1],  
       [ 8,  5],  
       [ 3,  7],  
       [ 3,  6],  
       [ 7,  3],  
       [ 1,  9],  
       [ 3,  8],  
       [ 2,  2],  
       [ 2,  1],  
       [10,  8]], dtype=int64)
```

```
y_train = y[0:15]  
y_train  
Out[17]:  
0      Fruit  
1    Protein  
2      Fruit  
3    Vegetable  
4    Vegetable  
5    Protein  
6      Fruit  
7    Vegetable  
8    Protein  
9      Fruit  
10   Vegetable  
11   Vegetable  
12    Protein  
13    Protein  
14      Fruit  
Name: Food Type, dtype: object  
  
#####  
X_test = X[15:16,]  
X_test  
Out[20]: array([[6, 4]], dtype=int64)
```

# Use 'Scikit-Learn' Package: kNeighborsClassifier Function

- k=1, Fruit (Fruit = 1)
- k=3, Fruit (Fruit = 2, Protein = 1)

```
clf = neighbors.KNeighborsClassifier(n_neighbors=1)
clf.fit(X_train, y_train)
Out[22]:
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=1, n_neighbors=1, p=2,
                    weights='uniform')

clf.predict(X_test)
Out[23]: array(['Fruit'], dtype=object)

#####
clf = neighbors.KNeighborsClassifier(n_neighbors=3)
clf.fit(X_train, y_train)
Out[26]:
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=1, n_neighbors=3, p=2,
                    weights='uniform')

clf.predict(X_test)
Out[27]: array(['Fruit'], dtype=object)
```

# Use 'Scikit-Learn' Package: kNeighborsClassifier Function

- k=5, Fruit or Protein (Fruit = 2, Protein = 2, Vegetable = 1)
- k=9, Protein (Fruit = 3, Protein = 4, Vegetable = 2)

```
clf = neighbors.KNeighborsClassifier(n_neighbors=5)
clf.fit(X_train, y_train)
Out[29]:
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=1, n_neighbors=5, p=2,
                    weights='uniform')

clf.predict(X_test)
Out[30]: array(['Fruit'], dtype=object)

#####
clf = neighbors.KNeighborsClassifier(n_neighbors=9)
clf.fit(X_train, y_train)
Out[33]:
KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
                    metric_params=None, n_jobs=1, n_neighbors=9, p=2,
                    weights='uniform')

clf.predict(X_test)
Out[34]: array(['Protein'], dtype=object)
```



# Summary

---

- Similarity Based Learning
- The Bayes Classifier
- 'k' Nearest Neighbor (kNN) Model
- kNN Model Assessment
- Data Normalization:  
Standardization & Scaling
- kNN in R
  - Example 1: Dataset = Food, No Package
  - Example 2: Dataset = Food, Package = Class
  - Example 3: Dataset = Iris, Package = Class
- kNN in Python
  - Example 4: Dataset = Food, No Package
  - Example 5: Dataset = Food, Package = Scikit-Learn